





Stochastic Gradient Estimation

Cheng Soon Ong
Marc Peter Deisenroth

December 2020



Monte Carlo Gradient Estimation in Machine Learning

Shakir Mohamed^{*1}

SHAKIR@GOOGLE.COM

Mihaela Rosca^{*1 2}

MIHAELACR@GOOGLE.COM

Michael Figurnov^{*1}

MFIGURNOV@GOOGLE.COM

Andriy Mnih^{*1}

AMNIH@GOOGLE.COM

**Equal contributions;*

1 DeepMind, London

2 University College London

Setting

- ▶ Expected utility

$$\mathbb{E}_{\boldsymbol{x} \sim p(\boldsymbol{x}; \boldsymbol{\theta})}[U(\boldsymbol{x})] = \int U(\boldsymbol{x})p(\boldsymbol{x}; \boldsymbol{\theta})d\boldsymbol{x}$$

- ▶ Distributional parameters $\boldsymbol{\theta}$

Setting

- ▶ Expected utility

$$\mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})] = \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x}$$

- ▶ Distributional parameters $\boldsymbol{\theta}$
- ▶ Gradients w.r.t. distributional parameters $\boldsymbol{\theta}$ of an expected utility:

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})]$$

Setting

- ▶ Expected utility

$$\mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})] = \int U(\mathbf{x})p(\mathbf{x}; \boldsymbol{\theta})d\mathbf{x}$$

- ▶ Distributional parameters $\boldsymbol{\theta}$
- ▶ Gradients w.r.t. distributional parameters $\boldsymbol{\theta}$ of an expected utility:

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})]$$

- ▶ Sensitivity analysis ➡ Explanation
- ▶ Training of machine learning models ➡ Optimization

Where?

- ▶ **Variational inference.** Gradient of evidence lower bound (ELBO) w.r.t. variational parameters θ :

$$\nabla_{\theta} \mathbb{E}_{z \sim q(z|x;\theta)} \left[\log p(x|z) - \log \frac{q(z|x;\theta)}{p(z)} \right]$$

Where?

- ▶ **Variational inference.** Gradient of evidence lower bound (ELBO) w.r.t. variational parameters θ :

$$\nabla_{\theta} \mathbb{E}_{z \sim q(z|x;\theta)} \left[\log p(x|z) - \log \frac{q(z|x;\theta)}{p(z)} \right]$$

- ▶ **Reinforcement learning.** Gradient of expected long-term reward w.r.t. policy parameters θ :

$$\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau;\theta)} \left[\sum_{t=0}^T \gamma^t r(\mathbf{x}_t, \mathbf{u}_t) \right], \quad \text{trajectory } \tau = (\mathbf{x}_0, \mathbf{u}_0, \dots, \mathbf{x}_T, \mathbf{u}_T)$$

Where?

- ▶ **Variational inference.** Gradient of evidence lower bound (ELBO) w.r.t. variational parameters θ :

$$\nabla_{\theta} \mathbb{E}_{z \sim q(z|x;\theta)} \left[\log p(x|z) - \log \frac{q(z|x;\theta)}{p(z)} \right]$$

- ▶ **Reinforcement learning.** Gradient of expected long-term reward w.r.t. policy parameters θ :

$$\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau;\theta)} \left[\sum_{t=0}^T \gamma^t r(\mathbf{x}_t, \mathbf{u}_t) \right], \quad \text{trajectory } \tau = (\mathbf{x}_0, \mathbf{u}_0, \dots, \mathbf{x}_T, \mathbf{u}_T)$$

- ▶ **Experimental design and Bayesian optimization.** Gradient of the probability of improvement w.r.t. designs θ (where to measure next?):

$$\nabla_{\theta} \mathbb{E}_{y \sim p(y;\theta)} [\mathbb{1}_{\{y > y_{\text{best}}\}}]$$

Gradients of expectations

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)} [U(x)]$$

- ▶ If we can compute (an approximation of) the expected value analytically (requires deterministic approximate inference), we can use the **chain rule** to get the gradient
- ▶ Otherwise (e.g., stochastic approximate inference, mini-batching, Monte-Carlo estimation), we need to compute **gradients of a stochastic estimator**

Stochastic gradient estimators

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \nabla_{\theta} \int U(x) p(x; \theta) dx$$

- ▶ Derivatives of measures: Directly differentiate the measure $p(x; \theta)$ w.r.t. θ
 - ▶ Score-function gradient estimators
 - ▶ Measure-valued gradient estimators

Stochastic gradient estimators

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \nabla_{\theta} \int U(x)p(x; \theta)dx$$

- ▶ Derivatives of measures: Directly differentiate the measure $p(x; \theta)$ w.r.t. θ
 - ▶ Score-function gradient estimators
 - ▶ Measure-valued gradient estimators
- ▶ Derivatives of paths: Differentiate through path the parameters θ take (via random variables x to U)
 - ▶ Pathwise gradient estimators

Stochastic gradient estimators

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \nabla_{\theta} \int U(x) p(x; \theta) dx$$

- ▶ Derivatives of measures: Directly differentiate the measure $p(x; \theta)$ w.r.t. θ
 - ▶ Score-function gradient estimators
 - ▶ Measure-valued gradient estimators
- ▶ Derivatives of paths: Differentiate through path the parameters θ take (via random variables x to U)
 - ▶ Pathwise gradient estimators
- ▶ Repeating pattern: Swap the order of differentiation and expectation
 - ▶▶ Gradients of deterministic quantities ✓
 - ▶▶ Monte-Carlo integration to compute the expectation ✓

Score-Function Gradient Estimators

Key insight

Key idea

Use log-derivative trick to turn the gradient of an expectation into an expectation of a gradient.

Key insight

Key idea

Use log-derivative trick to turn the gradient of an expectation into an expectation of a gradient.

► Log-derivative trick

$$\underbrace{\nabla_{\theta} \log p(\mathbf{x}; \boldsymbol{\theta})}_{\text{score}} = \frac{\nabla_{\theta} p(\mathbf{x}; \boldsymbol{\theta})}{p(\mathbf{x}; \boldsymbol{\theta})} \quad \Rightarrow \quad \nabla_{\theta} p(\mathbf{x}; \boldsymbol{\theta}) = p(\mathbf{x}; \boldsymbol{\theta}) \nabla_{\theta} \log p(\mathbf{x}; \boldsymbol{\theta})$$

Key insight

Key idea

Use log-derivative trick to turn the gradient of an expectation into an expectation of a gradient.

► Log-derivative trick

$$\underbrace{\nabla_{\theta} \log p(\mathbf{x}; \boldsymbol{\theta})}_{\text{score}} = \frac{\nabla_{\theta} p(\mathbf{x}; \boldsymbol{\theta})}{p(\mathbf{x}; \boldsymbol{\theta})} \quad \Rightarrow \quad \nabla_{\theta} p(\mathbf{x}; \boldsymbol{\theta}) = p(\mathbf{x}; \boldsymbol{\theta}) \nabla_{\theta} \log p(\mathbf{x}; \boldsymbol{\theta})$$

Score-function gradient estimator: Derivation

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})} [U(\mathbf{x})] = \nabla_{\boldsymbol{\theta}} \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x}$$

Expectation as
integration

Score-function gradient estimator: Derivation

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \mathbb{E}_{x \sim p(\mathbf{x}; \boldsymbol{\theta})} [U(\mathbf{x})] &= \nabla_{\boldsymbol{\theta}} \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x} \\ &= \int U(\mathbf{x}) \nabla_{\boldsymbol{\theta}} p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x}\end{aligned}$$

Expectation as
integration

Move gradient inside

Score-function gradient estimator: Derivation

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})] &= \nabla_{\boldsymbol{\theta}} \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x} && \text{Expectation as integration} \\ &= \int U(\mathbf{x}) \nabla_{\boldsymbol{\theta}} p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x} && \text{Move gradient inside} \\ &= \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) \nabla_{\boldsymbol{\theta}} \log p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x} && \text{Log-derivative trick}\end{aligned}$$

Score-function gradient estimator: Derivation

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] &= \nabla_{\theta} \int U(x)p(x; \theta)dx && \text{Expectation as integration} \\ &= \int U(x)\nabla_{\theta}p(x; \theta)dx && \text{Move gradient inside} \\ &= \int U(x)p(x; \theta)\nabla_{\theta} \log p(x; \theta)dx && \text{Log-derivative trick} \\ &= \mathbb{E}_{x \sim p(x; \theta)}[U(x)\nabla_{\theta} \log p(x; \theta)] && \text{Integral as expectation}\end{aligned}$$

Score-function gradient estimator

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)} [U(x)] = \mathbb{E}_{x \sim p(x; \theta)} [U(x) \underbrace{\nabla_{\theta} \log p(x; \theta)}_{\text{score}}]$$

- ▶ Turned gradient of an expectation into the expectation of a (deterministic) gradient
- ▶ Gradient is the expected utility-weighted score

Score-function gradient estimator

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \mathbb{E}_{x \sim p(x; \theta)}[U(x) \underbrace{\nabla_{\theta} \log p(x; \theta)}_{\text{score}}]$$

- ▶ Turned gradient of an expectation into the expectation of a (deterministic) gradient
- ▶ Gradient is the expected utility-weighted score
- ▶ Monte-Carlo integration to get the gradient:

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] \approx \frac{1}{S} \sum_{s=1}^S U(\mathbf{x}^{(s)}) \nabla_{\theta} \log p(\mathbf{x}^{(s)}; \theta), \quad \mathbf{x}^{(s)} \sim p(\mathbf{x}; \theta)$$

Properties: Score-function gradient estimator

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] \approx \frac{1}{S} \sum_{s=1}^S U(x^{(s)}) \nabla_{\theta} \log p(x^{(s)}; \theta), \quad x^{(s)} \sim p(x; \theta)$$

- ▶ Single-sample estimation is OK, i.e., $S = 1$
- ▶ Any type of utility function U can be used (e.g., non-differentiable)
- ▶ p must be differentiable w.r.t. θ
- ▶ Must be able to sample easily from $p(x; \theta)$
- ▶ Discrete and continuous distributions are OK
- ▶ Techniques to control the variance of the estimator
(e.g., Greensmith et al., 2004; Titsias & Lázaro-Gredilla, 2015)

Example: REINFORCE (Williams, 1992)

- ▶ Optimize parameters θ of a (stochastic) policy $p(\mathbf{u}_t | \mathbf{x}_t; \theta)$

$$\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau)]$$

- ▶ $U(\tau)$: Long-term reward for trajectory τ

Example: REINFORCE (Williams, 1992)

- ▶ Optimize parameters θ of a (stochastic) policy $p(\mathbf{u}_t | \mathbf{x}_t; \theta)$

$$\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau)]$$

- ▶ $U(\tau)$: Long-term reward for trajectory τ
- ▶ Trajectory distribution (with $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t, \epsilon_t)$)

$$p(\tau; \theta) = p(\mathbf{x}_0, \mathbf{u}_0, \dots, \mathbf{u}_T, \mathbf{x}_T; \theta) = p(\mathbf{x}_0) \prod_{t=0}^T \underbrace{p(\mathbf{x}_{t+1} | \mathbf{x}_t, \mathbf{u}_t)}_{\text{state transition}} \underbrace{p(\mathbf{u}_t | \mathbf{x}_t; \theta)}_{\text{policy}}$$

Example: REINFORCE (2) (Williams, 1992)

$$\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau)] = \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau) \nabla_{\theta} \log p(\tau; \theta)]$$

Example: REINFORCE (2) (Williams, 1992)

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau)] &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau) \nabla_{\theta} \log p(\tau; \theta)] \\ &= \mathbb{E}_{\tau \sim p(\tau; \theta)} \left[U(\tau) \nabla_{\theta} \left(\log p(x_0) + \sum_{t=0}^T \underbrace{\log p(x_{t+1} | x_t, u_t)}_{\text{state transition}} + \underbrace{\log p(u_t | x_t; \theta)}_{\text{policy}} \right) \right]\end{aligned}$$

Example: REINFORCE (2) (Williams, 1992)

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau)] &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau) \nabla_{\theta} \log p(\tau; \theta)] \\&= \mathbb{E}_{\tau \sim p(\tau; \theta)} \left[U(\tau) \nabla_{\theta} \left(\log p(x_0) + \sum_{t=0}^T \underbrace{\log p(x_{t+1} | x_t, u_t)}_{\text{state transition}} + \underbrace{\log p(u_t | x_t; \theta)}_{\text{policy}} \right) \right] \\&= \mathbb{E}_{\tau \sim p(\tau; \theta)} \left[U(\tau) \sum_{t=0}^T \nabla_{\theta} \log p(u_t | x_t; \theta) \right]\end{aligned}$$

Example: REINFORCE (2) (Williams, 1992)

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau)] &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [U(\tau) \nabla_{\theta} \log p(\tau; \theta)] \\&= \mathbb{E}_{\tau \sim p(\tau; \theta)} \left[U(\tau) \nabla_{\theta} \left(\log p(x_0) + \sum_{t=0}^T \underbrace{\log p(x_{t+1}|x_t, u_t)}_{\text{state transition}} + \underbrace{\log p(u_t|x_t; \theta)}_{\text{policy}} \right) \right] \\&= \mathbb{E}_{\tau \sim p(\tau; \theta)} \left[U(\tau) \sum_{t=0}^T \nabla_{\theta} \log p(u_t|x_t; \theta) \right]\end{aligned}$$

- ▶ Markov property of state evolution
 - ▶▶ Only need gradient of the log-policy at each time step
- ▶ Monte Carlo for expectation
- ▶ Can be used in model-free and model-based settings


Applications

- ▶ Reinforcement learning (e.g., Williams, 1992; Sutton et al., 2000)
- ▶ (Black-box) variational inference (e.g., Paisley et al., 2012, Ranganath et al., 2014)
- ▶ Discrete-event systems (operations research)
- ▶ Computational finance

Pathwise Gradient Estimators

Setting

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)]$$




$$x = f(z; \theta)$$

- ▶ Data x can be obtained by a **deterministic transformation f (path)** of a latent variable $z \sim p(z)$, where $p(z)$ has no tunable parameters, e.g., $p(z) = \mathcal{N}(\mathbf{0}, \mathbf{I})$
- ▶ Distributional parameters of $p(x; \theta)$ are the parameters of the path f
- ▶ **Push gradients through this path** (chain rule)

Setting

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)]$$




$$x = f(z; \theta)$$

- ▶ Data x can be obtained by a **deterministic transformation f (path)** of a latent variable $z \sim p(z)$, where $p(z)$ has no tunable parameters, e.g., $p(z) = \mathcal{N}(\mathbf{0}, \mathbf{I})$
- ▶ Distributional parameters of $p(x; \theta)$ are the parameters of the path f
- ▶ **Push gradients through this path** (chain rule)

Key idea

Define a path from a latent variable z to data x and use the change-of-variables trick to turn the gradient of an expectation into an expectation of a gradient.

Derivation




$$\mathbf{x} = f(\mathbf{z}; \boldsymbol{\theta})$$

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})} [U(\mathbf{x})] = \nabla_{\boldsymbol{\theta}} \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x}$$

Expectation as
integration

Derivation




$$\mathbf{x} = f(\mathbf{z}; \boldsymbol{\theta})$$

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})} [U(\mathbf{x})] &= \nabla_{\boldsymbol{\theta}} \int U(\mathbf{x}) p(\mathbf{x}; \boldsymbol{\theta}) d\mathbf{x} \\ &= \nabla_{\boldsymbol{\theta}} \int U(f(\mathbf{z}; \boldsymbol{\theta})) p(\mathbf{z}) d\mathbf{z}\end{aligned}$$

Expectation as
integration

Change of variables

Derivation



$$x = f(z; \theta)$$

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \nabla_{\theta} \int U(x)p(x; \theta)dx$$

Expectation as
integration


$$= \nabla_{\theta} \int U(f(z; \theta))p(z)dz$$

Change of variables

$$= \int \nabla_{\theta} U(f(z; \theta))p(z)dz$$

Move gradient inside

Derivation



$$x = f(z; \theta)$$

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] &= \nabla_{\theta} \int U(x) p(x; \theta) dx \\ &= \nabla_{\theta} \int U(f(z; \theta)) p(z) dz \\ &= \int \nabla_{\theta} U(f(z; \theta)) p(z) dz \\ &= \mathbb{E}_{z \sim p(z)}[\underbrace{\nabla_{\theta} U(f(z; \theta))}_{=x}]\end{aligned}$$


Expectation as
integration

Change of variables

Move gradient inside

Integral as expectation

Pathwise gradient estimator




$$x = f(z; \theta)$$

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \mathbb{E}_{z \sim p(z)}[\nabla_{\theta} U(f(z; \theta))]$$

- ▶ Turned gradient of an expectation into an expectation (w.r.t. a parameter-free distribution) of a (deterministic) gradient
 - ▶▶ Push parameters inside U

Pathwise gradient estimator




$$x = f(z; \theta)$$

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] = \mathbb{E}_{z \sim p(z)}[\nabla_{\theta} U(f(z; \theta))]$$

- ▶ Turned gradient of an expectation into an expectation (w.r.t. a parameter-free distribution) of a (deterministic) gradient
 - ▶▶ Push parameters inside U
- ▶ Monte-Carlo estimator of the gradient

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] \approx \frac{1}{S} \sum_{s=1}^S \nabla_{\theta} U(f(z^{(s)}; \theta)), \quad z^{(s)} \sim p(z)$$

Pathwise gradient estimator



$$\mathbf{x} = f(\mathbf{z}; \boldsymbol{\theta})$$

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})] = \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})}[\nabla_{\boldsymbol{\theta}} U(f(\mathbf{z}; \boldsymbol{\theta}))]$$

- ▶ Turned gradient of an expectation into an expectation (w.r.t. a parameter-free distribution) of a (deterministic) gradient
 - ▶▶ Push parameters inside U
- ▶ Monte-Carlo estimator of the gradient

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}; \boldsymbol{\theta})}[U(\mathbf{x})] \approx \frac{1}{S} \sum_{s=1}^S \nabla_{\boldsymbol{\theta}} U(f(\mathbf{z}^{(s)}; \boldsymbol{\theta})), \quad \mathbf{z}^{(s)} \sim p(\mathbf{z})$$

$$\nabla_{\boldsymbol{\theta}} U(f(\mathbf{z}; \boldsymbol{\theta})) = \nabla_{\mathbf{x}} U(\mathbf{x}) \nabla_{\boldsymbol{\theta}} f(\mathbf{z}; \boldsymbol{\theta}) \quad \text{▶▶ Chain rule}$$

Properties: Pathwise gradient estimator

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)}[U(x)] \approx \frac{1}{S} \sum_{s=1}^S \nabla_x U(\mathbf{x}^{(s)}) \nabla_{\theta} f(\mathbf{z}^{(s)}; \theta), \quad \mathbf{z}^{(s)} \sim p(\mathbf{z})$$

- ▶ Single-sample estimation OK, i.e., $S = 1$.
- ▶ Utility U must be differentiable
- ▶ Path f must be differentiable
- ▶ Need to be able to sample from $p(\mathbf{z})$, but not from $p(x; \theta)$
- ▶ Often lower variance than score-function gradient estimator
- ▶ Control variability of path f to control variance of the estimator

Example: Bayesian optimization (Wilson et al., 2018)

- ▶ Inner loop of Bayesian optimization: Where to measure next?
 - ▶▶ Maximize acquisition function
- ▶ Many acquisition functions can be written as expected utilities

$$\mathcal{L} = \mathbb{E}_{\mathbf{y} \sim p(\mathbf{y}; \boldsymbol{\theta})}[U(\mathbf{y})] = \int U(\mathbf{y})p(\mathbf{y}; \boldsymbol{\theta})d\mathbf{y}, \quad p(\mathbf{y}; \boldsymbol{\theta}) = \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

Example: Bayesian optimization (Wilson et al., 2018)

- ▶ Inner loop of Bayesian optimization: Where to measure next?
 - ▶▶ Maximize acquisition function
- ▶ Many acquisition functions can be written as expected utilities

$$\mathcal{L} = \mathbb{E}_{\mathbf{y} \sim p(\mathbf{y}; \boldsymbol{\theta})}[U(\mathbf{y})] = \int U(\mathbf{y})p(\mathbf{y}; \boldsymbol{\theta})d\mathbf{y}, \quad p(\mathbf{y}; \boldsymbol{\theta}) = \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

- ▶ Define path from $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ to \mathbf{y} ▶▶ Pathwise gradient estimation

Example: Bayesian optimization (Wilson et al., 2018)

Abbr.	Acquisition Function \mathcal{L}	Reparameterization	MM
EI	$\mathbb{E}_{\mathbf{y}}[\max(\text{ReLU}(\mathbf{y} - \alpha))]$	$\mathbb{E}_{\mathbf{z}}[\max(\text{ReLU}(\boldsymbol{\mu} + \mathbf{L}\mathbf{z} - \alpha))]$	Y
PI	$\mathbb{E}_{\mathbf{y}}[\max(\mathbb{1}^-(\mathbf{y} - \alpha))]$	$\mathbb{E}_{\mathbf{z}}[\max(\sigma(\frac{\boldsymbol{\mu} + \mathbf{L}\mathbf{z} - \alpha}{\tau}))]$	Y
SR	$\mathbb{E}_{\mathbf{y}}[\max(\mathbf{y})]$	$\mathbb{E}_{\mathbf{z}}[\max(\boldsymbol{\mu} + \mathbf{L}\mathbf{z})]$	Y
UCB	$\mathbb{E}_{\mathbf{y}}[\max(\boldsymbol{\mu} + \sqrt{\beta\pi/2} \boldsymbol{\gamma})]$	$\mathbb{E}_{\mathbf{z}}[\max(\boldsymbol{\mu} + \sqrt{\beta\pi/2} \mathbf{L}\mathbf{z})]$	Y
ES	$-\mathbb{E}_{\mathbf{y}_a}[H(\mathbb{E}_{\mathbf{y}_b \mathbf{y}_a}[\mathbb{1}^+(\mathbf{y}_b - \max(\mathbf{y}_b))])]$	$-\mathbb{E}_{\mathbf{z}_a}[H(\mathbb{E}_{\mathbf{z}_b}[\text{softmax}(\frac{\boldsymbol{\mu}_{b a} + \mathbf{L}_{b a}\mathbf{z}_b}{\tau})])]$	N
KG	$\mathbb{E}_{\mathbf{y}_a}[\max(\boldsymbol{\mu}_b + \boldsymbol{\Sigma}_{b,a}\boldsymbol{\Sigma}_{a,a}^{-1}(\mathbf{y}_a - \boldsymbol{\mu}_a))]$	$\mathbb{E}_{\mathbf{z}_a}[\max(\boldsymbol{\mu}_b + \boldsymbol{\Sigma}_{b,a}\boldsymbol{\Sigma}_{a,a}^{-1}\mathbf{L}_a\mathbf{z}_a)]$	N

Table 1: Examples of reparameterizable acquisition functions; the final column indicates whether they belong to the MM family (Section 3.2). Glossary: $\mathbb{1}^{+/-}$ denotes the right-/left-continuous Heaviside step function; ReLU and σ rectified linear and sigmoid nonlinearities, respectively; H the Shannon entropy; α an improvement threshold; τ a temperature parameter; $\mathbf{L}\mathbf{L}^\top \triangleq \boldsymbol{\Sigma}$ the Cholesky factor; and, residuals $\boldsymbol{\gamma} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma})$. Lastly, non-myopic acquisition function (ES and KG) are assumed to be defined using a discretization. Terms associated with the query set and discretization are respectively denoted via subscripts a and b .

From Wilson et al. (2018)

Application areas

- ▶ Bayesian optimization (e.g., Wilson et al., 2018)
- ▶ Normalizing flows (e.g., Rezende & Mohamed, 2015)
- ▶ Variational auto-encoders (e.g., Kingma & Welling, 2014; Rezende et al., 2014)
- ▶ Generative models (e.g., Goodfellow et al., 2014; Mohamed & Lakshminarayanan, 2016)
- ▶ Reinforcement learning (e.g., Heess et al., 2015)
- ▶ Probabilistic programming (e.g., Ritchie et al., 2016)

Summary

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x; \theta)} [U(x)]$$

- ▶ Compute gradient of an expected utility
- ▶ Key idea: Swap order of differentiation and integration (expectation)
 - ▶▶ Use Monte Carlo methods to compute gradients
- ▶ Score-function gradient estimator using log-derivative trick
- ▶ Pathwise gradient estimator defines a parametrized path from a latent variable to the data



References

- Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2014). Generative Adversarial Networks. In *Advances in Neural Information Processing Systems*.
- Greensmith, E., Bartlett, P. L., and Baxter, J. (2004). Variance Reduction Techniques for Gradient Estimates in Reinforcement Learning. *Journal of Machine Learning Research*, 5:1471–1530.
- Heess, N., Wayne, G., and Timothy Lillicrap, D. S., Tassa, Y., and Erez, T. (2015). Learning Continuous Control Policies by Stochastic Value Gradients. In *Advances in Neural Information Processing Systems*.
- Kingma, D. P. and Welling, M. (2014). Auto-Encoding Variational Bayes. In *Proceedings of the International Conference on Learning Representations*.
- Mohamed, S. and Lakshminarayanan, B. (2016). Learning in Implicit Generative Models. *arXiv:1610.03483*.
- Paisley, J., Blei, D. M., and Jordan, M. I. (2012). Variational Bayesian Inference with Stochastic Search. In *Proceedings of the International Conference on Machine Learning*.
- Ranganath, R., Gerrish, S., and Blei, D. M. (2014). Black Box Variational Inference. *arXiv:1401.0118*.

References (cont.)

- Rezende, D. J., Mohamed, S., and Wierstra, D. (2014). Stochastic Backpropagation and Variational Inference in Deep Latent Gaussian Models. In *Proceedings of the International Conference on Machine Learning*.
- Ritchie, D., Horsfall, P., and Goodman, N. D. (2016). Deep Amortized Inference for Probabilistic Programs. *arXiv:1610.05735*.
- Sutton, R. S., Mcallester, D. A., Singh, S. P., and Mansour, Y. (2000). Policy Gradient Methods for Reinforcement Learning with Function Approximation. In *Advances in Neural Information Processing Systems*.
- Titsias, M. K. and Lázaro-Gredilla, M. (2015). Local Expectation Gradients for Black Box Variational Inference. In *Advances in Neural Information Processing Systems*.
- Williams, R. J. (1992). Simple Statistical Gradient-following Algorithms for Connectionist Reinforcement Learning. *Machine Learning*, 8(3):229–256.
- Wilson, J. T., Hutter, F., and Deisenroth, M. P. (2018). Maximizing Acquisition Functions for Bayesian Optimization. In *Advances in Neural Information Processing Systems*.